

HW4

November 10, 2023

```
[828]: import pandas as pd
import numpy as np
import datasets
import torch
import math
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from sklearn.metrics import precision_score, recall_score, f1_score,
    ↪classification_report
from torch.utils.data import TensorDataset
import time
from itertools import chain
from torch.nn.utils.rnn import pad_sequence
import torch.nn.functional as F
import copy
from torch.optim import lr_scheduler
```

```
[48]: import torch
import math
# this ensures that the current MacOS version is at least 12.3+
print(torch.backends.mps.is_available())
# this ensures that the current PyTorch installation was built with MPS
    ↪activated.
print(torch.backends.mps.is_built())
```

True

True

```
[49]: dtype = torch.float
device = torch.device("mps")
```

```
[244]: dataset = datasets.load_dataset("conll2003")
```

0.0.1 Convert words/tokens to indices

```
[683]: import itertools
from collections import Counter

word_frequency = Counter(itertools.chain(*dataset['train']['tokens'])) # type: ignore

# Remove words below threshold 3
word_frequency = {
    word: frequency
    for word, frequency in word_frequency.items()
    if frequency >= 3
}

word2idx = {
    word: index
    for index, word in enumerate(word_frequency.keys(), start=2)
}

word2idx['[PAD]'] = 0
word2idx['[UNK]'] = 1
```

```
[1004]: sample_tokens = dataset['train'][0]['tokens']
sample_tokens
```

```
[1004]: ['EU', 'rejects', 'German', 'call', 'to', 'boycott', 'British', 'lamb', '.']
```

```
[684]: # the vocab size
vocab_size = max(word2idx.values())+1
vocab_size
```

```
[684]: 8128
```

```
[689]: def convert_word_to_id(sample):
    #Code to convert all tokens to their respective indexes
    #If the token is unknown, we set index of 1
    input_ids = [ word2idx.get(token, 1) for token in sample['tokens'] ]

    sample['input_ids'] = input_ids
    return sample

dataset = dataset.map(convert_word_to_id)
```

```
Map:   0%|          | 0/14041 [00:00<?, ? examples/s]
```

```
Map:   0%|          | 0/3250 [00:00<?, ? examples/s]
```

```
Map:   0%|          | 0/3453 [00:00<?, ? examples/s]
```

```
[692]: df_train = pd.DataFrame(dataset['train']).drop(columns=['pos_tags', 'chunk_tags', 'id', 'tokens'])
df_train.columns = ['label', 'input_ids']

df_test = pd.DataFrame(dataset['test']).drop(columns=['pos_tags', 'chunk_tags', 'id', 'tokens'])
df_test.columns = ['label', 'input_ids']

df_val = pd.DataFrame(dataset['validation']).drop(columns=['pos_tags', 'chunk_tags', 'id', 'tokens'])
df_val.columns = ['label', 'input_ids']
```

0.0.2 Padding

```
[1145]: import pandas as pd
import torch
from torch.utils.data import Dataset

# Create a custom Dataset class
class CustomDataset(Dataset):
    def __init__(self, dataframe):
        self.data = dataframe

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        label = torch.tensor(self.data.loc[idx, "label"], dtype=torch.long)
        input_ids = torch.tensor(self.data.loc[idx, "input_ids"], dtype=torch.
        ↪long)

        return input_ids, label

# Create an instance of the CustomDataset
dataset_train = CustomDataset(df_train)

# Example: Accessing a single sample
print(dataset_train[2])
```

```
(tensor([12, 13]), tensor([5, 0]))
```

```
[1146]: def custom_collate(batch):
    # Separate input sequences and labels
    input_seqs, labels = zip(*batch)

    # Calculate the sequence lengths based on input sequences (assuming they
    ↪have the same length as labels)
```

```

sequence_lengths = [len(seq) for seq in input_seqs]

# Sort input sequences and labels by sequence length (descending)
sorted_seqs_and_labels = sorted(zip(input_seqs, labels), key=lambda x:
↪len(x[0]), reverse=True)
sorted_input_seqs, sorted_labels = zip(*sorted_seqs_and_labels)

# Pad input sequences to the maximum length within the batch
padded_input_seqs = pad_sequence(sorted_input_seqs, batch_first=True,
↪padding_value=0) # Use 0 as the padding value
padded_labels = pad_sequence(sorted_labels, batch_first=True,
↪padding_value=0) # Use 0 as the padding value

return padded_input_seqs, padded_labels

```

0.0.3 Create dataloaders

```

[1147]: def dataloader_generator(df,shuffle):
        dataset_from_df = CustomDataset(df)
        batch_size = 64
        dataloader = DataLoader(dataset_from_df, batch_size=batch_size,
↪collate_fn=custom_collate, shuffle=shuffle)
        return dataloader

train_loader = dataloader_generator(df_train,shuffle=True)
test_loader = dataloader_generator(df_test,shuffle=False)
val_loader = dataloader_generator(df_val,shuffle=False)

for batch in val_loader:
    input_val, target_val = batch
    break

```

0.0.4 Building the model

```

[340]: !wget https://raw.githubusercontent.com/sighsmile/conlleva/master/conlleva.py

```

```

--2023-11-06 16:54:57--
https://raw.githubusercontent.com/sighsmile/conlleva/master/conlleva.py
  raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133,
185.199.109.133, 185.199.110.133, ...
  raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443...
  HTTP ... 200 OK
  7502 (7.3K) [text/plain]
  : "conlleva.py.1"

```

```

conlleva.py.1      100%[=====>]    7.33K  --.-KB/s    0s

```

```
[341]: from conlleval import evaluate
```

```
[1150]: class BiLSTMNER(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim,
        ↪ num_layers, dropout):
        super(BiLSTMNER, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.bilstm = nn.LSTM(embedding_dim, hidden_dim, num_layers=num_layers,
            batch_first=True, bidirectional=True)
        self.dropout = nn.Dropout(dropout)
        self.linear = nn.Linear(hidden_dim * 2, output_dim)
        self.elu = nn.ELU()
        self.classifier = nn.Linear(output_dim, num_tags) # num_tags is the
        ↪ number of unique NER tags

    def forward(self, x):
        x = self.embedding(x)
        x, _ = self.bilstm(x)
        x = self.dropout(x)
        x = self.linear(x)
        x = self.elu(x)
        x = self.classifier(x)
        return x

#initialize
num_tags = 9
vocab_size = max(word2idx.values())+1

model = BiLSTMNER(vocab_size, 100, 256, 128, 1, 0.33)
optimizer = optim.Adam(model.parameters(), lr=0.001)
loss_function = nn.CrossEntropyLoss()

#training
num_epochs = 20
print('start training')
for epoch in range(num_epochs):
    start_time = time.time()
    model.train()
    total_loss = 0
    for batch in train_loader:
        optimizer.zero_grad()
        inputs, targets = batch
        outputs = model(inputs)
```

```

        batch_size = inputs.size()[-1]
        #From the instruction of CrossEntropy, we need to change the format of
        ↪outputs
        loss = loss_function(outputs.permute(0,2,1), targets)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    end_time = time.time()
    print(f'Epoch {epoch + 1}, Loss: {total_loss / len(train_loader)}, time:
    ↪{end_time-start_time}s')
    print('validation error: ')
    precision, recall, f1 = eval(model, val_loader)

```

start training

Epoch 1, Loss: 0.24208735396916217, time: 52.45967507362366s

validation error:

processed 152266 tokens with 5942 phrases; found: 2102 phrases; correct: 1080.

accuracy: 20.82%; (non-0)

accuracy: 95.34%; precision: 51.38%; recall: 18.18%; FB1: 26.85

LOC: precision: 59.36%; recall: 26.08%; FB1: 36.23 807

MISC: precision: 33.33%; recall: 0.33%; FB1: 0.64 9

ORG: precision: 30.00%; recall: 3.36%; FB1: 6.04 150

PER: precision: 48.68%; recall: 30.02%; FB1: 37.14 1136

Epoch 2, Loss: 0.11488994293930856, time: 47.83418798446655s

validation error:

processed 152266 tokens with 5942 phrases; found: 4661 phrases; correct: 2929.

accuracy: 52.38%; (non-0)

accuracy: 97.05%; precision: 62.84%; recall: 49.29%; FB1: 55.25

LOC: precision: 72.27%; recall: 65.11%; FB1: 68.50 1655

MISC: precision: 62.36%; recall: 37.20%; FB1: 46.60 550

ORG: precision: 45.43%; recall: 44.44%; FB1: 44.93 1312

PER: precision: 69.41%; recall: 43.11%; FB1: 53.18 1144

Epoch 3, Loss: 0.06882393922318111, time: 48.614689111709595s

validation error:

processed 152266 tokens with 5942 phrases; found: 5102 phrases; correct: 3706.

accuracy: 66.22%; (non-0)

accuracy: 97.86%; precision: 72.64%; recall: 62.37%; FB1: 67.11

LOC: precision: 85.42%; recall: 69.84%; FB1: 76.85 1502

MISC: precision: 65.60%; recall: 57.70%; FB1: 61.40 811

ORG: precision: 62.67%; recall: 54.21%; FB1: 58.14 1160

PER: precision: 71.45%; recall: 63.19%; FB1: 67.07 1629

Epoch 4, Loss: 0.04578833337026564, time: 51.33672094345093s

validation error:

processed 152266 tokens with 5942 phrases; found: 5280 phrases; correct: 4048.

accuracy: 71.87%; (non-0)

accuracy: 98.20%; precision: 76.67%; recall: 68.13%; FB1: 72.14

LOC: precision: 89.86%; recall: 74.80%; FB1: 81.64 1529

MISC: precision: 75.29%; recall: 63.45%; FB1: 68.86 777
 ORG: precision: 62.88%; recall: 63.16%; FB1: 63.02 1347
 PER: precision: 76.34%; recall: 67.43%; FB1: 71.61 1627
 Epoch 5, Loss: 0.032878815653649245, time: 48.04787993431091s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5393 phrases; correct: 4290.
 accuracy: 75.31%; (non-0)
 accuracy: 98.40%; precision: 79.55%; recall: 72.20%; FB1: 75.69
 LOC: precision: 85.29%; recall: 80.19%; FB1: 82.66 1727
 MISC: precision: 80.28%; recall: 68.87%; FB1: 74.14 791
 ORG: precision: 73.72%; recall: 64.21%; FB1: 68.63 1168
 PER: precision: 77.39%; recall: 71.72%; FB1: 74.44 1707
 Epoch 6, Loss: 0.023866635279475964, time: 48.24751901626587s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5465 phrases; correct: 4370.
 accuracy: 76.53%; (non-0)
 accuracy: 98.46%; precision: 79.96%; recall: 73.54%; FB1: 76.62
 LOC: precision: 88.65%; recall: 79.91%; FB1: 84.05 1656
 MISC: precision: 80.98%; recall: 69.74%; FB1: 74.94 794
 ORG: precision: 71.36%; recall: 67.26%; FB1: 69.25 1264
 PER: precision: 77.50%; recall: 73.67%; FB1: 75.54 1751
 Epoch 7, Loss: 0.017713668648238208, time: 48.44618272781372s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 6117 phrases; correct: 4592.
 accuracy: 80.59%; (non-0)
 accuracy: 98.36%; precision: 75.07%; recall: 77.28%; FB1: 76.16
 LOC: precision: 85.71%; recall: 82.63%; FB1: 84.15 1771
 MISC: precision: 78.41%; recall: 70.50%; FB1: 74.24 829
 ORG: precision: 62.20%; recall: 70.92%; FB1: 66.27 1529
 PER: precision: 74.09%; recall: 79.97%; FB1: 76.92 1988
 Epoch 8, Loss: 0.013745928631926125, time: 47.203505992889404s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5289 phrases; correct: 4409.
 accuracy: 76.44%; (non-0)
 accuracy: 98.53%; precision: 83.36%; recall: 74.20%; FB1: 78.51
 LOC: precision: 91.49%; recall: 80.78%; FB1: 85.81 1622
 MISC: precision: 80.94%; recall: 72.78%; FB1: 76.64 829
 ORG: precision: 78.06%; recall: 66.07%; FB1: 71.57 1135
 PER: precision: 80.33%; recall: 74.27%; FB1: 77.18 1703
 Epoch 9, Loss: 0.010567720067179338, time: 47.350847005844116s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5375 phrases; correct: 4429.
 accuracy: 76.74%; (non-0)
 accuracy: 98.51%; precision: 82.40%; recall: 74.54%; FB1: 78.27
 LOC: precision: 91.60%; recall: 81.27%; FB1: 86.13 1630
 MISC: precision: 80.24%; recall: 73.54%; FB1: 76.74 845
 ORG: precision: 71.33%; recall: 69.20%; FB1: 70.25 1301
 PER: precision: 83.18%; recall: 72.20%; FB1: 77.30 1599

Epoch 10, Loss: 0.00821540692069737, time: 46.872527837753296s
validation error:
processed 152266 tokens with 5942 phrases; found: 5612 phrases; correct: 4468.
accuracy: 77.35%; (non-0)
accuracy: 98.49%; precision: 79.62%; recall: 75.19%; FB1: 77.34
LOC: precision: 84.22%; recall: 84.81%; FB1: 84.51 1850
MISC: precision: 81.68%; recall: 71.58%; FB1: 76.30 808
ORG: precision: 71.56%; recall: 67.56%; FB1: 69.51 1266
PER: precision: 79.62%; recall: 72.96%; FB1: 76.15 1688

Epoch 11, Loss: 0.006528246736640788, time: 46.133893966674805s
validation error:
processed 152266 tokens with 5942 phrases; found: 6042 phrases; correct: 4622.
accuracy: 80.33%; (non-0)
accuracy: 98.43%; precision: 76.50%; recall: 77.79%; FB1: 77.14
LOC: precision: 86.12%; recall: 83.78%; FB1: 84.93 1787
MISC: precision: 71.69%; recall: 75.27%; FB1: 73.44 968
ORG: precision: 71.24%; recall: 68.53%; FB1: 69.86 1290
PER: precision: 73.61%; recall: 79.80%; FB1: 76.58 1997

Epoch 12, Loss: 0.00525556694959629, time: 48.15459370613098s
validation error:
processed 152266 tokens with 5942 phrases; found: 5796 phrases; correct: 4560.
accuracy: 79.07%; (non-0)
accuracy: 98.49%; precision: 78.67%; recall: 76.74%; FB1: 77.70
LOC: precision: 86.47%; recall: 83.12%; FB1: 84.76 1766
MISC: precision: 79.93%; recall: 73.43%; FB1: 76.54 847
ORG: precision: 68.11%; recall: 70.40%; FB1: 69.23 1386
PER: precision: 78.58%; recall: 76.66%; FB1: 77.60 1797

Epoch 13, Loss: 0.00421205094052394, time: 47.617199182510376s
validation error:
processed 152266 tokens with 5942 phrases; found: 5846 phrases; correct: 4574.
accuracy: 79.32%; (non-0)
accuracy: 98.51%; precision: 78.24%; recall: 76.98%; FB1: 77.60
LOC: precision: 86.80%; recall: 83.02%; FB1: 84.86 1757
MISC: precision: 80.05%; recall: 73.54%; FB1: 76.65 847
ORG: precision: 67.91%; recall: 70.40%; FB1: 69.13 1390
PER: precision: 77.05%; recall: 77.47%; FB1: 77.26 1852

Epoch 14, Loss: 0.0038494242876450616, time: 47.95925307273865s
validation error:
processed 152266 tokens with 5942 phrases; found: 5703 phrases; correct: 4530.
accuracy: 78.86%; (non-0)
accuracy: 98.53%; precision: 79.43%; recall: 76.24%; FB1: 77.80
LOC: precision: 87.94%; recall: 82.96%; FB1: 85.38 1733
MISC: precision: 77.65%; recall: 75.38%; FB1: 76.50 895
ORG: precision: 73.76%; recall: 67.71%; FB1: 70.61 1231
PER: precision: 76.08%; recall: 76.17%; FB1: 76.13 1844

Epoch 15, Loss: 0.0037216038500032895, time: 49.66430187225342s
validation error:
processed 152266 tokens with 5942 phrases; found: 5539 phrases; correct: 4432.

accuracy: 77.23%; (non-0)
 accuracy: 98.50%; precision: 80.01%; recall: 74.59%; FB1: 77.21
 LOC: precision: 88.67%; recall: 82.25%; FB1: 85.34 1704
 MISC: precision: 73.08%; recall: 74.19%; FB1: 73.63 936
 ORG: precision: 73.55%; recall: 67.19%; FB1: 70.23 1225
 PER: precision: 79.81%; recall: 72.53%; FB1: 76.00 1674
 Epoch 16, Loss: 0.003172349494839595, time: 49.49557089805603s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5817 phrases; correct: 4554.
 accuracy: 78.83%; (non-0)
 accuracy: 98.49%; precision: 78.29%; recall: 76.64%; FB1: 77.46
 LOC: precision: 85.91%; recall: 83.61%; FB1: 84.74 1788
 MISC: precision: 74.62%; recall: 74.30%; FB1: 74.46 918
 ORG: precision: 71.13%; recall: 68.53%; FB1: 69.81 1292
 PER: precision: 77.74%; recall: 76.76%; FB1: 77.25 1819
 Epoch 17, Loss: 0.002620459050971972, time: 47.30130100250244s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5804 phrases; correct: 4577.
 accuracy: 79.24%; (non-0)
 accuracy: 98.50%; precision: 78.86%; recall: 77.03%; FB1: 77.93
 LOC: precision: 87.66%; recall: 83.51%; FB1: 85.53 1750
 MISC: precision: 76.78%; recall: 73.86%; FB1: 75.29 887
 ORG: precision: 69.38%; recall: 70.47%; FB1: 69.92 1362
 PER: precision: 78.50%; recall: 76.93%; FB1: 77.71 1805
 Epoch 18, Loss: 0.00238024852177742, time: 47.947713136672974s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5805 phrases; correct: 4541.
 accuracy: 78.55%; (non-0)
 accuracy: 98.47%; precision: 78.23%; recall: 76.42%; FB1: 77.31
 LOC: precision: 85.42%; recall: 84.21%; FB1: 84.81 1811
 MISC: precision: 76.91%; recall: 72.99%; FB1: 74.90 875
 ORG: precision: 69.67%; recall: 70.25%; FB1: 69.96 1352
 PER: precision: 78.04%; recall: 74.86%; FB1: 76.42 1767
 Epoch 19, Loss: 0.002281301094626542, time: 48.1134819984436s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5658 phrases; correct: 4538.
 accuracy: 78.69%; (non-0)
 accuracy: 98.56%; precision: 80.21%; recall: 76.37%; FB1: 78.24
 LOC: precision: 87.74%; recall: 84.16%; FB1: 85.91 1762
 MISC: precision: 77.63%; recall: 73.75%; FB1: 75.64 876
 ORG: precision: 75.79%; recall: 67.93%; FB1: 71.65 1202
 PER: precision: 77.06%; recall: 76.06%; FB1: 76.56 1818
 Epoch 20, Loss: 0.002566172640349991, time: 48.22484111785889s
 validation error:
 processed 152266 tokens with 5942 phrases; found: 5651 phrases; correct: 4545.
 accuracy: 78.65%; (non-0)
 accuracy: 98.55%; precision: 80.43%; recall: 76.49%; FB1: 78.41
 LOC: precision: 88.78%; recall: 83.56%; FB1: 86.09 1729

```

MISC: precision: 78.97%; recall: 74.95%; FB1: 76.91 875
ORG: precision: 74.96%; recall: 67.64%; FB1: 71.11 1210
PER: precision: 76.86%; recall: 76.66%; FB1: 76.76 1837

```

```
[1162]: print(f"Validation: precision = {precision}, recall = {recall}, f1 = {f1}")
```

```

precision = 80.42824278888693, recall = 76.48939750925614, f1 =
78.40938497369102

```

```
[1151]: # SAVE THE MODEL
torch.save(model.state_dict(), 'task1.pth')
```

```
[444]: # Example reversed_ner_tags dictionary
reversed_ner_tags = {
    0: 'O',
    1: 'B-PER',
    2: 'I-PER',
    3: 'B-ORG',
    4: 'I-ORG',
    5: 'B-LOC',
    6: 'I-LOC',
    7: 'B-MISC',
    8: 'I-MISC'
}

# Example tensor with shape (32, 36)
tensor = torch.randint(0, 9, (32, 36)) # Random integers between 0 and 8

# Map tensor elements using reversed_ner_tags
mapped_tensor = [[reversed_ner_tags[item.item()] for item in row] for row in
↪ tensor]
```

```
[398]: ner_tags = {'O': 0, 'B-PER': 1, 'I-PER': 2, 'B-ORG': 3, 'I-ORG': 4, 'B-LOC': 5,
↪ 'I-LOC': 6, 'B-MISC': 7, 'I-MISC': 8}

reversed_ner_tags = {value: key for key, value in ner_tags.items()}
reversed_ner_tags
```

```
[398]: {0: 'O',
1: 'B-PER',
2: 'I-PER',
3: 'B-ORG',
4: 'I-ORG',
5: 'B-LOC',
6: 'I-LOC',
7: 'B-MISC',
8: 'I-MISC'}
```

```
[1149]: #evaluation
def eval(model, loader):
    model.eval()
    all_preds, all_labels = [], []
    with torch.no_grad():
        for batch in loader:
            inputs, targets = batch
            outputs = model(inputs)
            _, preds = torch.max(outputs, -1)
            preds_converted = [[reversed_ner_tags[item.item()]] for item in row]
    for row in preds:
        targets_converted = [[reversed_ner_tags[item.item()]] for item in
    row] for row in targets:
        all_preds.extend(preds_converted)
        all_labels.extend(targets_converted)
    # all_preds = list(chain.from_iterable(all_preds))
    # all_labels = list(chain.from_iterable(all_labels))
    # all_labels = torch.cat(all_labels)
    all_preds = itertools.chain(*all_preds)
    all_labels = itertools.chain(*all_labels)
    result = evaluate(all_labels, all_preds, verbose=True)
    precision, recall, f1 = result[0], result[1], result[2]
    return precision, recall, f1
```

```
[1163]: print('Test: ')
precision, recall, f1 = eval(model, test_loader)
```

Test:

processed 146937 tokens with 5648 phrases; found: 5146 phrases; correct: 3710.

accuracy: 70.02%; (non-0)

accuracy: 97.95%; precision: 72.09%; recall: 65.69%; FB1: 68.74

LOC: precision: 84.52%; recall: 75.30%; FB1: 79.64 1486

MISC: precision: 64.47%; recall: 62.82%; FB1: 63.64 684

ORG: precision: 67.13%; recall: 57.80%; FB1: 62.12 1430

PER: precision: 68.11%; recall: 65.12%; FB1: 66.58 1546

```
[1164]: print(f"Test: precision = {precision}, recall = {recall}, f1 = {f1}")
```

Test: precision = 72.09483093664983, recall = 65.68696883852692, f1 = 68.74189364461739

0.0.5 Solution for the task 1

1. Hyperparameters:

- vocab_size = 8128
- embedding_dim = 100
- hidden_dim = 256
- output_dim = 128

- num_layers = 1
 - dropout = 0.33
 - optimizer learning rate= 0.001
 - batch_size = 64
2. Solution: At first, I created a vocab that maps all the tokens from the training set to a number, and I gave up the tokens that appeared less than 3 times. Secondly, I customized a dataset class so that each batch will contain (input_ids, ner_tags). Next, I used padding_sequence to customize the padding value of 0 in input and 9 in ner_tags. Why do I pad here? I need to make sure for each batch, which contains 32 samples, will have the max_length within one batch. Thirdly, I designed my bilstm model. The model will firstly embed all the inputs to 100-dim vectors and then throw the vectors to the lstm layer. Through elu, dropout, and one more linear layer, it model will predict the name entity for each token in samples.
3. Questions and answers:
- What are the precision, recall, and F1 score on the validation data?
 - precision = 80.42824278888693, recall = 76.48939750925614, f1 = 78.40938497369102- What are the precision, recall, and F1 score on the test data?
 - precision = 72.09483093664983, recall = 65.68696883852692, f1 = 68.74189364461739

0.1 Task 2: Glove Embedding

0.1.1 Load Glove Embedding

```
[31]: # Define a function to load GloVe embeddings from a file
def load_glove_embeddings(file_path):
    embeddings_index = {}
    with open(file_path, encoding="utf-8") as f:
        for line in f:
            values = line.split()
            word = values[0]
            coefs = np.asarray(values[1:], dtype="float32")
            embeddings_index[word] = coefs
    return embeddings_index

# Specify the path to your downloaded "glove.6B.100d.txt" file
glove_file_path = "glove.6B.100d"

# Load GloVe embeddings into memory
glove_embeddings = load_glove_embeddings(glove_file_path)
```

0.1.2 Create Glove Idx

```
[1027]: def convert_word_to_glove_ids(sample):
    tokens = sample['tokens']
    glove_ids = []
    for token in tokens:
        token = token.lower()
        indices = np.where(vocab_npa == token)
```

```

        if indices[0].size > 0:
            index = indices[0][0]
        else:
            index = 1
        glove_ids.append(index)
        sample['glove_ids'] = glove_ids
    return sample
dataset = dataset.map(convert_word_to_glove_ids)

```

Map: 0%| | 0/14041 [00:00<?, ? examples/s]

Map: 0%| | 0/3250 [00:00<?, ? examples/s]

Map: 0%| | 0/3453 [00:00<?, ? examples/s]

0.1.3 Customize the layer

```

[831]: #convert glove into a layer
vocab, embeddings = [], []
with open('glove.6B.100d', encoding="utf-8") as fi:
    full_content = fi.read().strip().split('\n')
for i in range(len(full_content)):
    i_word = full_content[i].split(' ')[0]
    i_embeddings = [float(val) for val in full_content[i].split(' ')[1:]]
    vocab.append(i_word)
    embeddings.append(i_embeddings)

```

```

[832]: vocab_npa = np.array(vocab)
embs_npa = np.array(embeddings)

#insert '<pad>' and '<unk>' tokens at start of vocab_npa.
vocab_npa = np.insert(vocab_npa, 0, '<pad>')
vocab_npa = np.insert(vocab_npa, 1, '<unk>')
print(vocab_npa[:10])

pad_emb_npa = np.zeros((1, embs_npa.shape[1])) #embedding for '<pad>' token.
unk_emb_npa = np.mean(embs_npa, axis=0, keepdims=True) #embedding for '<unk>'
↳ token.

#insert embeddings for pad and unk tokens at top of embs_npa.
embs_npa = np.vstack((pad_emb_npa, unk_emb_npa, embs_npa))
print(embs_npa.shape)

```

```

['<pad>' '<unk>' 'the' ', ' '.' 'of' 'to' 'and' 'in' 'a']
(400002, 100)

```

```

[877]: import torch
my_embedding_layer = torch.nn.Embedding.from_pretrained(torch.
↳ from_numpy(embs_npa).float(), freeze=True)

```

```
assert my_embedding_layer.weight.shape == embs_npa.shape
print(my_embedding_layer.weight.shape)
```

```
torch.Size([400002, 100])
```

0.1.4 Make Glove case-sensitive – creating another feature

```
[858]: #add features to the dataloader
#case 0: lower case - no uppercase
#case 1: first word is uppercase
#case 2: whole word is uppercase
#case 3: others: e.g. ", "
def capital_case(word):
    if word.islower():
        return 0
    elif word.isupper():
        return 2
    elif word.istitle():
        return 1
    else: return 3

def convert_word_to_capital_case(sample):
    capitals = [capital_case(word) for word in sample['tokens']]
    sample['capitals'] =capitals
    return sample

dataset = dataset.map(convert_word_to_capital_case)
```

```
Map:   0%|          | 0/14041 [00:00<?, ? examples/s]
```

```
Map:   0%|          | 0/3250 [00:00<?, ? examples/s]
```

```
Map:   0%|          | 0/3453 [00:00<?, ? examples/s]
```

```
[1023]: dataset['train'][2]
```

```
[1023]: {'id': '2',
'tokens': ['BRUSSELS', '1996-08-22'],
'pos_tags': [22, 11],
'chunk_tags': [11, 12],
'ner_tags': [5, 0],
'input_ids': [12, 13],
'capitals': [2, 3],
'glove_ids': [1, 1]}
```

0.1.5 Padding – glove embedding

```
[1165]: import pandas as pd
import torch
from torch.utils.data import Dataset

# Create a custom Dataset class
class CustomDataset(Dataset):

    def __init__(self, data):
        self.data = data

    def __len__(self):
        return len(self.data)

    def __getitem__(self, index):
        label = torch.tensor(self.data[index]['ner_tags'], dtype=torch.long )
        glove_ids = torch.tensor(self.data[index]['glove_ids'], dtype=torch.
↪long)
        capital = torch.tensor(self.data[index]['capitals'], dtype=torch.long)

        return label, glove_ids, capital

# Create an instance of the CustomDataset
dataset_train = CustomDataset(dataset['train'])
dataset_test = CustomDataset(dataset['test'])
dataset_val = CustomDataset(dataset['validation'])

# Example: Accessing a single sample
print(dataset_train[0])

(tensor([3, 0, 7, 0, 0, 0, 7, 0, 0]), tensor([ 646, 7580, 516, 582, 6,
5262, 299, 10240, 4]), tensor([2, 0, 1, 0, 0, 0, 1, 0, 3]))
```

```
[1166]: def custom_collate(batch):
    label, glove_ids, capital = zip(*batch)
    padded_label = pad_sequence(label, batch_first=True, padding_value=9 )
    padded_glove_ids = pad_sequence(glove_ids, batch_first=True, ↪
↪padding_value=0 )
    padded_capital = pad_sequence(capital, batch_first=True, padding_value=4 )
    return padded_glove_ids, padded_capital, padded_label
```

```
[1167]: batch_size = 64
train_loader = DataLoader(dataset_train, batch_size=batch_size, collate_fn=↪
↪custom_collate, shuffle=True)
test_loader = DataLoader(dataset_test, batch_size=batch_size, collate_fn=↪
↪custom_collate, shuffle=False)
```

```
val_loader = DataLoader(dataset_val, batch_size=batch_size, collate_fn=
↳ custom_collate, shuffle=False)
```

```
[1174]: class BiLSTMNER(nn.Module):
    def __init__(self, hidden_dim, output_dim, num_layers, dropout):
        super(BiLSTMNER, self).__init__()
        self.embedding = my_embedding_layer
        self.capital_layer = nn.
↳ Embedding(num_embeddings=5, embedding_dim=20, padding_idx=4)
        self.bilstm = nn.LSTM(input_size=120, hidden_size=hidden_dim,
↳ num_layers=num_layers,
                                batch_first=True, bidirectional=True)
        self.dropout = nn.Dropout(dropout)
        self.linear = nn.Linear(hidden_dim * 2, output_dim, dtype=torch.float32)
        self.elu = nn.ELU()
        self.classifier = nn.Linear(output_dim, num_tags, dtype=torch.float32)
↳ # num_tags is the number of unique NER tags

    def forward(self, x, capital):
        x = self.embedding(x.int())
        capital = self.capital_layer(capital.int())
        x = torch.cat([x, capital], dim=2)
        x, _ = self.bilstm(x)
        x = self.dropout(x)
        x = self.linear(x)
        x = self.elu(x)
        x = self.classifier(x)
        return x

#initialize
num_tags = 9
model = BiLSTMNER(256, 128, 1, 0.33)
optimizer = optim.Adam(model.parameters(), lr=0.001)
loss_function = nn.CrossEntropyLoss(ignore_index=9)

#training
num_epochs = 20
print('start training')
for epoch in range(num_epochs):
    start_time = time.time()
    model.train()
    total_loss = 0
    for batch in train_loader:
        optimizer.zero_grad()
        inputs, capitals, targets = batch
```



```

        outputs = model(inputs, capitals)
        batch_size = inputs.size()[-1]
        #From the instruction of CrossEntropy, we need to change the format of
        ↪outputs
        loss = loss_function(outputs.permute(0,2,1), targets)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
        end_time = time.time()
        print(f'Epoch {epoch + 1}, Loss: {total_loss / len(train_loader)}, time:␣
        ↪{end_time-start_time}s')
        print('validation error: ')
        precision, recall, f1 = eval(model, val_loader)

```

start training

Epoch 1, Loss: 0.2787948575378819, time: 59.194642066955566s

validation error:

processed 51362 tokens with 5942 phrases; found: 6183 phrases; correct: 5010.

accuracy: 85.81%; (non-0)

accuracy: 97.16%; precision: 81.03%; recall: 84.32%; FB1: 82.64

LOC: precision: 83.66%; recall: 90.85%; FB1: 87.11 1995

MISC: precision: 68.96%; recall: 74.95%; FB1: 71.83 1002

ORG: precision: 72.70%; recall: 70.69%; FB1: 71.68 1304

PER: precision: 90.44%; recall: 92.40%; FB1: 91.41 1882

Epoch 2, Loss: 0.0857482789422978, time: 59.39733099937439s

validation error:

processed 51362 tokens with 5942 phrases; found: 6033 phrases; correct: 5285.

accuracy: 89.63%; (non-0)

accuracy: 97.99%; precision: 87.60%; recall: 88.94%; FB1: 88.27

LOC: precision: 92.77%; recall: 92.16%; FB1: 92.46 1825

MISC: precision: 79.46%; recall: 79.28%; FB1: 79.37 920

ORG: precision: 78.05%; recall: 84.04%; FB1: 80.93 1444

PER: precision: 94.03%; recall: 94.14%; FB1: 94.09 1844

Epoch 3, Loss: 0.06427998816255819, time: 58.05035185813904s

validation error:

processed 51362 tokens with 5942 phrases; found: 6028 phrases; correct: 5366.

accuracy: 91.06%; (non-0)

accuracy: 98.25%; precision: 89.02%; recall: 90.31%; FB1: 89.66

LOC: precision: 94.65%; recall: 92.49%; FB1: 93.56 1795

MISC: precision: 80.36%; recall: 82.54%; FB1: 81.43 947

ORG: precision: 81.43%; recall: 85.98%; FB1: 83.64 1416

PER: precision: 93.74%; recall: 95.17%; FB1: 94.45 1870

Epoch 4, Loss: 0.05143080727959221, time: 58.354299783706665s

validation error:

processed 51362 tokens with 5942 phrases; found: 6084 phrases; correct: 5425.

accuracy: 92.00%; (non-0)

accuracy: 98.33%; precision: 89.17%; recall: 91.30%; FB1: 90.22

LOC: precision: 93.02%; recall: 95.75%; FB1: 94.37 1891
 MISC: precision: 78.62%; recall: 83.73%; FB1: 81.09 982
 ORG: precision: 84.84%; recall: 83.89%; FB1: 84.36 1326
 PER: precision: 93.85%; recall: 96.04%; FB1: 94.93 1885
 Epoch 5, Loss: 0.04187326981178061, time: 59.04758620262146s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6045 phrases; correct: 5450.
 accuracy: 92.28%; (non-0)
 accuracy: 98.47%; precision: 90.16%; recall: 91.72%; FB1: 90.93
 LOC: precision: 93.03%; recall: 95.16%; FB1: 94.08 1879
 MISC: precision: 83.41%; recall: 82.86%; FB1: 83.13 916
 ORG: precision: 84.62%; recall: 87.40%; FB1: 85.99 1385
 PER: precision: 94.69%; recall: 95.87%; FB1: 95.28 1865
 Epoch 6, Loss: 0.03367272468114441, time: 58.262818813323975s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6143 phrases; correct: 5483.
 accuracy: 93.15%; (non-0)
 accuracy: 98.44%; precision: 89.26%; recall: 92.28%; FB1: 90.74
 LOC: precision: 94.36%; recall: 95.59%; FB1: 94.97 1861
 MISC: precision: 79.81%; recall: 82.75%; FB1: 81.26 956
 ORG: precision: 82.64%; recall: 89.49%; FB1: 85.93 1452
 PER: precision: 94.13%; recall: 95.77%; FB1: 94.94 1874
 Epoch 7, Loss: 0.027345544093457814, time: 59.46216917037964s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6015 phrases; correct: 5484.
 accuracy: 92.49%; (non-0)
 accuracy: 98.54%; precision: 91.17%; recall: 92.29%; FB1: 91.73
 LOC: precision: 94.70%; recall: 95.32%; FB1: 95.01 1849
 MISC: precision: 82.86%; recall: 86.01%; FB1: 84.41 957
 ORG: precision: 87.71%; recall: 87.84%; FB1: 87.78 1343
 PER: precision: 94.43%; recall: 95.66%; FB1: 95.04 1866
 Epoch 8, Loss: 0.021476747551721267, time: 58.77496004104614s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6044 phrases; correct: 5502.
 accuracy: 92.69%; (non-0)
 accuracy: 98.55%; precision: 91.03%; recall: 92.60%; FB1: 91.81
 LOC: precision: 94.57%; recall: 95.75%; FB1: 95.16 1860
 MISC: precision: 85.90%; recall: 83.95%; FB1: 84.91 901
 ORG: precision: 84.22%; recall: 89.93%; FB1: 86.98 1432
 PER: precision: 95.25%; recall: 95.71%; FB1: 95.48 1851
 Epoch 9, Loss: 0.01826574724400416, time: 60.954275131225586s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6081 phrases; correct: 5504.
 accuracy: 93.13%; (non-0)
 accuracy: 98.58%; precision: 90.51%; recall: 92.63%; FB1: 91.56
 LOC: precision: 95.46%; recall: 94.99%; FB1: 95.23 1828
 MISC: precision: 82.56%; recall: 84.71%; FB1: 83.62 946
 ORG: precision: 84.06%; recall: 90.83%; FB1: 87.31 1449

PER: precision: 94.73%; recall: 95.55%; FB1: 95.14 1858
Epoch 10, Loss: 0.013936886461239986, time: 64.49255204200745s
validation error:
processed 51362 tokens with 5942 phrases; found: 6052 phrases; correct: 5498.
accuracy: 92.89%; (non-0)
accuracy: 98.58%; precision: 90.85%; recall: 92.53%; FB1: 91.68
LOC: precision: 95.55%; recall: 94.67%; FB1: 95.11 1820
MISC: precision: 81.27%; recall: 87.53%; FB1: 84.28 993
ORG: precision: 86.99%; recall: 87.77%; FB1: 87.38 1353
PER: precision: 94.11%; recall: 96.36%; FB1: 95.23 1886
Epoch 11, Loss: 0.010753243909725412, time: 63.68161940574646s
validation error:
processed 51362 tokens with 5942 phrases; found: 6062 phrases; correct: 5484.
accuracy: 92.71%; (non-0)
accuracy: 98.53%; precision: 90.47%; recall: 92.29%; FB1: 91.37
LOC: precision: 94.12%; recall: 95.92%; FB1: 95.01 1872
MISC: precision: 83.39%; recall: 86.01%; FB1: 84.68 951
ORG: precision: 84.46%; recall: 87.55%; FB1: 85.98 1390
PER: precision: 94.92%; recall: 95.28%; FB1: 95.10 1849
Epoch 12, Loss: 0.009369476515249433, time: 61.54601192474365s
validation error:
processed 51362 tokens with 5942 phrases; found: 6029 phrases; correct: 5496.
accuracy: 92.99%; (non-0)
accuracy: 98.62%; precision: 91.16%; recall: 92.49%; FB1: 91.82
LOC: precision: 95.07%; recall: 95.43%; FB1: 95.25 1844
MISC: precision: 81.38%; recall: 86.77%; FB1: 83.99 983
ORG: precision: 87.74%; recall: 87.02%; FB1: 87.38 1330
PER: precision: 94.87%; recall: 96.42%; FB1: 95.64 1872
Epoch 13, Loss: 0.007336435311431573, time: 62.67510199546814s
validation error:
processed 51362 tokens with 5942 phrases; found: 6031 phrases; correct: 5527.
accuracy: 93.33%; (non-0)
accuracy: 98.66%; precision: 91.64%; recall: 93.02%; FB1: 92.32
LOC: precision: 94.47%; recall: 96.79%; FB1: 95.62 1882
MISC: precision: 83.98%; recall: 84.71%; FB1: 84.34 930
ORG: precision: 87.96%; recall: 89.86%; FB1: 88.90 1370
PER: precision: 95.35%; recall: 95.71%; FB1: 95.53 1849
Epoch 14, Loss: 0.006210239743284712, time: 63.491557121276855s
validation error:
processed 51362 tokens with 5942 phrases; found: 6091 phrases; correct: 5548.
accuracy: 93.79%; (non-0)
accuracy: 98.69%; precision: 91.09%; recall: 93.37%; FB1: 92.21
LOC: precision: 95.59%; recall: 95.65%; FB1: 95.62 1838
MISC: precision: 85.22%; recall: 86.33%; FB1: 85.78 934
ORG: precision: 85.04%; recall: 91.13%; FB1: 87.98 1437
PER: precision: 94.21%; recall: 96.25%; FB1: 95.22 1882
Epoch 15, Loss: 0.0048910202573593286, time: 89.17146420478821s
validation error:

processed 51362 tokens with 5942 phrases; found: 6044 phrases; correct: 5511.
 accuracy: 93.11%; (non-0)
 accuracy: 98.61%; precision: 91.18%; recall: 92.75%; FB1: 91.96
 LOC: precision: 94.52%; recall: 95.86%; FB1: 95.19 1863
 MISC: precision: 86.20%; recall: 86.01%; FB1: 86.10 920
 ORG: precision: 87.30%; recall: 87.62%; FB1: 87.46 1346
 PER: precision: 93.05%; recall: 96.74%; FB1: 94.86 1915
 Epoch 16, Loss: 0.00643773535636931, time: 87.57236385345459s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6053 phrases; correct: 5531.
 accuracy: 93.32%; (non-0)
 accuracy: 98.64%; precision: 91.38%; recall: 93.08%; FB1: 92.22
 LOC: precision: 94.51%; recall: 96.46%; FB1: 95.47 1875
 MISC: precision: 84.08%; recall: 85.36%; FB1: 84.71 936
 ORG: precision: 88.36%; recall: 88.29%; FB1: 88.33 1340
 PER: precision: 94.01%; recall: 97.07%; FB1: 95.51 1902
 Epoch 17, Loss: 0.004882407614059048, time: 63.116442918777466s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6020 phrases; correct: 5503.
 accuracy: 92.97%; (non-0)
 accuracy: 98.60%; precision: 91.41%; recall: 92.61%; FB1: 92.01
 LOC: precision: 94.92%; recall: 95.70%; FB1: 95.31 1852
 MISC: precision: 85.82%; recall: 84.71%; FB1: 85.26 910
 ORG: precision: 86.98%; recall: 89.19%; FB1: 88.07 1375
 PER: precision: 93.89%; recall: 95.98%; FB1: 94.93 1883
 Epoch 18, Loss: 0.003682805795291312, time: 133.74659514427185s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6058 phrases; correct: 5526.
 accuracy: 93.25%; (non-0)
 accuracy: 98.63%; precision: 91.22%; recall: 93.00%; FB1: 92.10
 LOC: precision: 94.57%; recall: 96.62%; FB1: 95.58 1877
 MISC: precision: 84.57%; recall: 86.23%; FB1: 85.39 940
 ORG: precision: 87.42%; recall: 88.07%; FB1: 87.74 1351
 PER: precision: 93.92%; recall: 96.36%; FB1: 95.12 1890
 Epoch 19, Loss: 0.003912460297711236, time: 155.75212383270264s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6046 phrases; correct: 5518.
 accuracy: 93.14%; (non-0)
 accuracy: 98.63%; precision: 91.27%; recall: 92.86%; FB1: 92.06
 LOC: precision: 94.93%; recall: 95.81%; FB1: 95.37 1854
 MISC: precision: 83.76%; recall: 86.12%; FB1: 84.92 948
 ORG: precision: 88.31%; recall: 87.92%; FB1: 88.12 1335
 PER: precision: 93.50%; recall: 96.91%; FB1: 95.17 1909
 Epoch 20, Loss: 0.0027471507286959836, time: 69.84025812149048s
 validation error:
 processed 51362 tokens with 5942 phrases; found: 6044 phrases; correct: 5541.
 accuracy: 93.47%; (non-0)
 accuracy: 98.67%; precision: 91.68%; recall: 93.25%; FB1: 92.46

```

LOC: precision: 94.80%; recall: 96.30%; FB1: 95.54 1866
MISC: precision: 85.01%; recall: 86.12%; FB1: 85.56 934
ORG: precision: 88.16%; recall: 89.41%; FB1: 88.78 1360
PER: precision: 94.43%; recall: 96.58%; FB1: 95.49 1884

```

```

[1179]: # SAVE THE MODEL
torch.save(model.state_dict(), 'task2.pth')

```

```

[1176]: print(f"Validation: precision = {precision}, recall = {recall}, f1 = {f1}")

Validation: precision = 91.67769688947716, recall = 93.25143049478291, f1 =
92.45786751209747

```

```

[1172]: ner_tags = {'O': 0, 'B-PER': 1, 'I-PER': 2, 'B-ORG': 3, 'I-ORG': 4, 'B-LOC': 5,
↳ 'I-LOC': 6, 'B-MISC': 7, 'I-MISC': 8, '<PAD>': 9}

reversed_ner_tags = {value: key for key, value in ner_tags.items()}
reversed_ner_tags

```

```

[1172]: {0: 'O',
1: 'B-PER',
2: 'I-PER',
3: 'B-ORG',
4: 'I-ORG',
5: 'B-LOC',
6: 'I-LOC',
7: 'B-MISC',
8: 'I-MISC',
9: '<PAD>'}

```

```

[1177]: #evaluation
def eval(model, loader):
    model.eval()
    all_preds, all_labels = [], []
    with torch.no_grad():
        for batch in loader:
            inputs, capitals, targets = batch
            #get rid of paddings on targets
            label_unpad = targets
            mask = label_unpad != 9
            label_unpad = label_unpad[mask]

            outputs = model(inputs, capitals)
            _, preds = torch.max(outputs, -1)
            #get rid of paddings on pred
            preds = preds[mask]

            preds_converted = [reversed_ner_tags[elem.item()] for elem in preds]

```

```

        targets_converted = [reversed_ner_tags[elem.item()] for elem in
↪label_unpad]
        all_preds.extend(preds_converted)
        all_labels.extend(targets_converted)
        # all_preds = list(chain.from_iterable(all_preds))
        # all_labels = list(chain.from_iterable(all_labels))
        # all_labels = torch.cat(all_labels)
        # all_preds = itertools.chain(*all_preds)
        # all_labels = itertools.chain(*all_labels)
        result = evaluate(all_labels, all_preds, verbose=True)
        precision, recall, f1 = result[0], result[1], result[2]
        return precision, recall, f1

# print('Test: ')
# precision, recall, f1 = eval(model, test_loader)

```

```
[1178]: print(f"Test: precision = {precision}, recall = {recall}, f1 = {f1}")
```

```
Test: precision = 91.67769688947716, recall = 93.25143049478291, f1 =
92.45786751209747
```

0.1.6 Solution for task2

1. Hyperparameters:

- embedding_dim = 100
- hidden_dim = 256
- output_dim = 128
- num_layers = 1
- dropout = 0.33
- optimizer learning rate= 0.001
- ignore_index = 9
- batch_size = 64

2. Solution: At first, I loaded the glove embedding and convert it into two arrays. One records all the indices and the other one records the 100-d embeddings for all the tokens. Secondly, since the glove is not case-sensitive, I tried to divide tokens into 4 cases (0: lowercase 1: some uppercases 2: all uppercases 3: lowercase and uppercase are the same). So, I added a new list to the dataset. Thirdly, I mapped all the tokens into indices in the glove embedding. So, I added one more list to the dataset. Forthly, I created a new customized dataset that each batch contains (glove_ids, capitalize, ner_tag). And similiar to the task, I padded all of them while creating the dataloaders. To be notified, I padded 9 to the ner_tag since it is a number that has not been used. I padded the capitalize with 4, which is not used either. Fifthly, I threw the batches into the model, which has the similar structure to the task 1. However, I added one more embedding layer such that the feature capitalize will be converted into 20-d vector and be added to the original 100-d layer. So, the input will become a 120-d vector. Through elu, dropout, and one more linear layer, it model will predict the name entity for each token in samples.

3. Questions and answers:

- What is the precision, recall, and F1 score on the validation data?
- precision = 91.67769688947716, recall = 93.25143049478291, f1 = 92.45786751209747- What are the precision, recall, and F1 score on the test data?
- precision = 91.67769688947716, recall = 93.25143049478291, f1 = 92.45786751209747- BiLSTM with Glove Embeddings outperforms the model without. Can you provide a rationale for this?
- At first, the glove is a bigger vocab than the word2idx, so it will map less unknown words. Secondly, since I added a new embedding layer, the model can better capture whether the word has been capitalized.

0.2 Task3: Transformer

```
[1180]: import pandas as pd
import torch
from torch.utils.data import Dataset

class CustomDataset(Dataset):

    def __init__(self, data):
        self.data = data

    def __len__(self):
        return len(self.data)

    def __getitem__(self, index):
        input = torch.tensor(self.data[index]['input_ids'], dtype=torch.long)
        target = torch.tensor(self.data[index]['ner_tags'], dtype=torch.long)

        return input, target

# Create an instance of the CustomDataset
dataset_train = CustomDataset(dataset['train'])
dataset_test = CustomDataset(dataset['test'])
dataset_val = CustomDataset(dataset['validation'])

# Example: Accessing a single sample
print(dataset_train[2])
```

(tensor([12, 13]), tensor([5, 0]))

```
[1181]: def custom_collate(batch):
    input_ids, label = zip(*batch)
    padded_label = pad_sequence(label, batch_first=True, padding_value=9)
    padded_input_ids = pad_sequence(input_ids, batch_first=True,
    ↪padding_value=0)
    return padded_input_ids, padded_label
```

```
[1182]: batch_size = 32
train_loader = DataLoader(dataset_train, batch_size=batch_size, collate_fn=
    ↪custom_collate, shuffle=True)
test_loader = DataLoader(dataset_test, batch_size=batch_size, collate_fn=
    ↪custom_collate, shuffle=False)
val_loader = DataLoader(dataset_val, batch_size=batch_size, collate_fn=
    ↪custom_collate, shuffle=False)

[1183]: # inspect the trainloader
for batch in train_loader:
    inputs, labels = batch
    break

[1184]: import torch
import torch.nn as nn

class TransformerNERModel(nn.Module):
    def __init__(self, vocab_size, tag_vocab_size, embed_size=128, num_heads=8,
    ↪max_seq_length=128, ff_dim=128, num_encoder_layers=6,
        dropout=0.33):
        super(TransformerNERModel, self).__init__()

        # Token embedding layer
        self.embedding = TokenEmbedding(vocab_size, embed_size)

        # Positional encoding
        self.positional_encoder = PositionalEncoding(emb_size= embed_size,
    ↪maxlen=max_seq_length)

        # Transformer Encoder
        self.transformer_encoder = nn.TransformerEncoder(
            nn.TransformerEncoderLayer(
                d_model=embed_size,
                nhead=num_heads,
                dim_feedforward=ff_dim,
                batch_first=True
            ),
            num_layers=num_encoder_layers,
        )

        # Linear layer for classification
        self.fc = nn.Linear(embed_size, tag_vocab_size)

    def forward(self, src, src_padding_mask):
        # Token embedding
        x = self.embedding(src)
```



```

        # Add positional encoding
        x = self.positional_encoder(x)

        # Transformer encoder
        x = self.transformer_encoder(x, src_key_padding_mask=src_padding_mask)

        # Final linear layer for classification
        x = self.fc(x)

    return x

class PositionalEncoding(nn.Module):
    def __init__(self,
                  emb_size: int,
                  dropout: float = 0.33,
                  maxlen: int = 5000):
        super(PositionalEncoding, self).__init__()
        den = torch.exp(- torch.arange(0, emb_size, 2)* math.log(10000) /
↪emb_size)
        pos = torch.arange(0, maxlen).reshape(maxlen, 1)
        pos_embedding = torch.zeros((maxlen, emb_size))
        pos_embedding[:, 0::2] = torch.sin(pos * den)
        pos_embedding[:, 1::2] = torch.cos(pos * den)
        pos_embedding = pos_embedding.unsqueeze(-2)

        self.dropout = nn.Dropout(dropout)
        self.register_buffer('pos_embedding', pos_embedding)

    def forward(self, token_embedding):
        return self.dropout(token_embedding + self.pos_embedding[:
↪token_embedding.size(0), :])

class TokenEmbedding(nn.Module):
    def __init__(self, vocab_size: int, emb_size):
        super(TokenEmbedding, self).__init__()
        self.embedding = nn.Embedding(vocab_size, emb_size)
        self.emb_size = emb_size

    def forward(self, tokens):
        return self.embedding(tokens.long()) * math.sqrt(self.emb_size)

# Initialize the model
vocab_size = max(word2idx.values())+1 # Your vocabulary size
tag_vocab_size = 9 # Your tag vocabulary size
model = TransformerNERModel(vocab_size, tag_vocab_size)

criterion = nn.CrossEntropyLoss(ignore_index=9)

```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
[1187]: # Training loop
num_epochs = 25
for epoch in range(num_epochs):
    model.train()
    total_loss = 0.0

    # Iterate over your training data in batches
    for batch_idx, (inputs, targets) in enumerate(train_loader):
        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass + src_padding_mask
        src_padding_mask = (inputs == 0).float()
        outputs = model(inputs, src_padding_mask= src_padding_mask)

        # Flatten the outputs and targets for the loss calculation
        outputs = outputs.view(-1, 9)
        targets = targets.view(-1)

        # Calculate the loss
        loss = criterion(outputs, targets)

        # Backpropagation
        loss.backward()
        optimizer.step()

        total_loss += loss.item()

    # Print the average loss for this epoch
    avg_loss = total_loss / len(train_loader)
    print(f"Epoch [{epoch + 1}/{num_epochs}] - Loss: {avg_loss:.4f}")

    print('validation error: ')
    precision, recall, f1 = eval(model, val_loader)
```

Epoch [1/20] - Loss: 0.5256

validation error:

processed 51362 tokens with 5942 phrases; found: 3827 phrases; correct: 1857.

accuracy: 29.40%; (non-0)

accuracy: 87.05%; precision: 48.52%; recall: 31.25%; FB1: 38.02

LOC: precision: 58.47%; recall: 52.04%; FB1: 55.07 1635

MISC: precision: 64.33%; recall: 23.86%; FB1: 34.81 342

ORG: precision: 53.95%; recall: 17.30%; FB1: 26.20 430

PER: precision: 31.62%; recall: 24.38%; FB1: 27.53 1420

Epoch [2/20] - Loss: 0.4477

validation error:

processed 51362 tokens with 5942 phrases; found: 3723 phrases; correct: 2158.
 accuracy: 34.28%; (non-0)
 accuracy: 88.53%; precision: 57.96%; recall: 36.32%; FB1: 44.66
 LOC: precision: 75.40%; recall: 54.38%; FB1: 63.19 1325
 MISC: precision: 72.88%; recall: 41.97%; FB1: 53.27 531
 ORG: precision: 48.84%; recall: 29.83%; FB1: 37.04 819
 PER: precision: 35.50%; recall: 20.20%; FB1: 25.74 1048
 Epoch [3/20] - Loss: 0.3974
 validation error:
 processed 51362 tokens with 5942 phrases; found: 4277 phrases; correct: 2363.
 accuracy: 39.20%; (non-0)
 accuracy: 89.35%; precision: 55.25%; recall: 39.77%; FB1: 46.25
 LOC: precision: 80.41%; recall: 57.43%; FB1: 67.01 1312
 MISC: precision: 67.03%; recall: 53.15%; FB1: 59.29 731
 ORG: precision: 51.54%; recall: 31.25%; FB1: 38.90 813
 PER: precision: 28.08%; recall: 21.66%; FB1: 24.46 1421
 Epoch [4/20] - Loss: 0.3614
 validation error:
 processed 51362 tokens with 5942 phrases; found: 5168 phrases; correct: 2721.
 accuracy: 44.73%; (non-0)
 accuracy: 89.73%; precision: 52.65%; recall: 45.79%; FB1: 48.98
 LOC: precision: 78.47%; recall: 61.89%; FB1: 69.20 1449
 MISC: precision: 67.09%; recall: 57.92%; FB1: 62.17 796
 ORG: precision: 55.21%; recall: 34.75%; FB1: 42.65 844
 PER: precision: 28.09%; recall: 31.70%; FB1: 29.79 2079
 Epoch [5/20] - Loss: 0.3321
 validation error:
 processed 51362 tokens with 5942 phrases; found: 5498 phrases; correct: 2887.
 accuracy: 49.01%; (non-0)
 accuracy: 90.45%; precision: 52.51%; recall: 48.59%; FB1: 50.47
 LOC: precision: 82.54%; recall: 61.51%; FB1: 70.49 1369
 MISC: precision: 69.62%; recall: 57.92%; FB1: 63.23 767
 ORG: precision: 45.41%; recall: 47.58%; FB1: 46.47 1405
 PER: precision: 29.89%; recall: 31.76%; FB1: 30.80 1957
 Epoch [6/20] - Loss: 0.3095
 validation error:
 processed 51362 tokens with 5942 phrases; found: 5344 phrases; correct: 2915.
 accuracy: 49.69%; (non-0)
 accuracy: 90.79%; precision: 54.55%; recall: 49.06%; FB1: 51.66
 LOC: precision: 76.19%; recall: 65.16%; FB1: 70.25 1571
 MISC: precision: 66.55%; recall: 61.50%; FB1: 63.92 852
 ORG: precision: 50.60%; recall: 37.96%; FB1: 43.37 1006
 PER: precision: 33.52%; recall: 34.85%; FB1: 34.18 1915
 Epoch [7/20] - Loss: 0.2901
 validation error:
 processed 51362 tokens with 5942 phrases; found: 5740 phrases; correct: 2861.
 accuracy: 49.97%; (non-0)
 accuracy: 90.56%; precision: 49.84%; recall: 48.15%; FB1: 48.98

LOC: precision: 77.28%; recall: 66.09%; FB1: 71.24 1571
 MISC: precision: 64.77%; recall: 61.61%; FB1: 63.15 877
 ORG: precision: 50.54%; recall: 31.17%; FB1: 38.56 827
 PER: precision: 26.82%; recall: 35.88%; FB1: 30.69 2465

Epoch [8/20] - Loss: 0.2772

validation error:

processed 51362 tokens with 5942 phrases; found: 5131 phrases; correct: 3051.

accuracy: 51.10%; (non-0)

accuracy: 91.32%; precision: 59.46%; recall: 51.35%; FB1: 55.11

LOC: precision: 82.48%; recall: 66.14%; FB1: 73.41 1473
 MISC: precision: 68.18%; recall: 63.45%; FB1: 65.73 858
 ORG: precision: 54.11%; recall: 48.10%; FB1: 50.93 1192
 PER: precision: 37.69%; recall: 32.90%; FB1: 35.13 1608

Epoch [9/20] - Loss: 0.2623

validation error:

processed 51362 tokens with 5942 phrases; found: 5718 phrases; correct: 3226.

accuracy: 54.99%; (non-0)

accuracy: 91.36%; precision: 56.42%; recall: 54.29%; FB1: 55.33

LOC: precision: 85.86%; recall: 62.17%; FB1: 72.12 1330
 MISC: precision: 71.46%; recall: 64.10%; FB1: 67.58 827
 ORG: precision: 49.14%; recall: 53.09%; FB1: 51.04 1449
 PER: precision: 36.98%; recall: 42.40%; FB1: 39.50 2112

Epoch [10/20] - Loss: 0.2519

validation error:

processed 51362 tokens with 5942 phrases; found: 5262 phrases; correct: 3086.

accuracy: 52.47%; (non-0)

accuracy: 91.59%; precision: 58.65%; recall: 51.94%; FB1: 55.09

LOC: precision: 83.68%; recall: 65.05%; FB1: 73.20 1428
 MISC: precision: 72.64%; recall: 66.81%; FB1: 69.60 848
 ORG: precision: 54.62%; recall: 52.42%; FB1: 53.50 1287
 PER: precision: 33.67%; recall: 31.05%; FB1: 32.31 1699

Epoch [11/20] - Loss: 0.2380

validation error:

processed 51362 tokens with 5942 phrases; found: 5797 phrases; correct: 3229.

accuracy: 54.00%; (non-0)

accuracy: 91.25%; precision: 55.70%; recall: 54.34%; FB1: 55.01

LOC: precision: 81.90%; recall: 66.03%; FB1: 73.12 1481
 MISC: precision: 70.90%; recall: 66.59%; FB1: 68.68 866
 ORG: precision: 50.79%; recall: 52.42%; FB1: 51.60 1384
 PER: precision: 33.83%; recall: 37.95%; FB1: 35.77 2066

Epoch [12/20] - Loss: 0.2313

validation error:

processed 51362 tokens with 5942 phrases; found: 5443 phrases; correct: 3330.

accuracy: 54.59%; (non-0)

accuracy: 91.61%; precision: 61.18%; recall: 56.04%; FB1: 58.50

LOC: precision: 81.57%; recall: 68.43%; FB1: 74.42 1541
 MISC: precision: 72.21%; recall: 65.94%; FB1: 68.93 842
 ORG: precision: 54.06%; recall: 54.06%; FB1: 54.06 1341

PER: precision: 43.05%; recall: 40.17%; FB1: 41.56 1719

Epoch [13/20] - Loss: 0.2247

validation error:

processed 51362 tokens with 5942 phrases; found: 5448 phrases; correct: 3308.

accuracy: 55.54%; (non-0)

accuracy: 91.77%; precision: 60.72%; recall: 55.67%; FB1: 58.09

LOC: precision: 80.15%; recall: 68.37%; FB1: 73.80 1567

MISC: precision: 75.03%; recall: 69.41%; FB1: 72.11 853

ORG: precision: 55.00%; recall: 45.56%; FB1: 49.84 1111

PER: precision: 41.78%; recall: 43.49%; FB1: 42.62 1917

Epoch [14/20] - Loss: 0.2153

validation error:

processed 51362 tokens with 5942 phrases; found: 5521 phrases; correct: 3226.

accuracy: 55.07%; (non-0)

accuracy: 91.79%; precision: 58.43%; recall: 54.29%; FB1: 56.29

LOC: precision: 79.81%; recall: 68.86%; FB1: 73.93 1585

MISC: precision: 70.17%; recall: 68.11%; FB1: 69.12 895

ORG: precision: 54.26%; recall: 46.53%; FB1: 50.10 1150

PER: precision: 37.49%; recall: 38.49%; FB1: 37.99 1891

Epoch [15/20] - Loss: 0.2109

validation error:

processed 51362 tokens with 5942 phrases; found: 5244 phrases; correct: 3283.

accuracy: 54.82%; (non-0)

accuracy: 91.86%; precision: 62.60%; recall: 55.25%; FB1: 58.70

LOC: precision: 79.68%; recall: 68.32%; FB1: 73.56 1575

MISC: precision: 71.86%; recall: 67.03%; FB1: 69.36 860

ORG: precision: 55.08%; recall: 50.11%; FB1: 52.48 1220

PER: precision: 46.44%; recall: 40.07%; FB1: 43.02 1589

Epoch [16/20] - Loss: 0.2058

validation error:

processed 51362 tokens with 5942 phrases; found: 5285 phrases; correct: 3317.

accuracy: 55.93%; (non-0)

accuracy: 92.03%; precision: 62.76%; recall: 55.82%; FB1: 59.09

LOC: precision: 82.97%; recall: 68.43%; FB1: 75.00 1515

MISC: precision: 72.47%; recall: 69.09%; FB1: 70.74 879

ORG: precision: 57.44%; recall: 49.81%; FB1: 53.35 1163

PER: precision: 43.69%; recall: 40.99%; FB1: 42.30 1728

Epoch [17/20] - Loss: 0.2001

validation error:

processed 51362 tokens with 5942 phrases; found: 5265 phrases; correct: 3308.

accuracy: 55.90%; (non-0)

accuracy: 91.97%; precision: 62.83%; recall: 55.67%; FB1: 59.03

LOC: precision: 84.92%; recall: 65.92%; FB1: 74.23 1426

MISC: precision: 68.81%; recall: 71.80%; FB1: 70.28 962

ORG: precision: 55.74%; recall: 50.71%; FB1: 53.10 1220

PER: precision: 45.56%; recall: 40.99%; FB1: 43.16 1657

Epoch [18/20] - Loss: 0.1967

validation error:

```

processed 51362 tokens with 5942 phrases; found: 5334 phrases; correct: 3472.
accuracy: 57.34%; (non-0)
accuracy: 92.18%; precision: 65.09%; recall: 58.43%; FB1: 61.58
    LOC: precision: 83.97%; recall: 67.88%; FB1: 75.08 1485
    MISC: precision: 75.44%; recall: 69.96%; FB1: 72.59 855
    ORG: precision: 56.91%; recall: 56.23%; FB1: 56.56 1325
    PER: precision: 49.49%; recall: 44.84%; FB1: 47.05 1669
Epoch [19/20] - Loss: 0.1907
validation error:
processed 51362 tokens with 5942 phrases; found: 5652 phrases; correct: 3506.
accuracy: 57.48%; (non-0)
accuracy: 91.84%; precision: 62.03%; recall: 59.00%; FB1: 60.48
    LOC: precision: 79.36%; recall: 71.15%; FB1: 75.03 1647
    MISC: precision: 78.36%; recall: 70.28%; FB1: 74.10 827
    ORG: precision: 57.29%; recall: 50.41%; FB1: 53.63 1180
    PER: precision: 43.79%; recall: 47.50%; FB1: 45.57 1998
Epoch [20/20] - Loss: 0.1868
validation error:
processed 51362 tokens with 5942 phrases; found: 5710 phrases; correct: 3461.
accuracy: 58.29%; (non-0)
accuracy: 92.03%; precision: 60.61%; recall: 58.25%; FB1: 59.41
    LOC: precision: 83.07%; recall: 70.50%; FB1: 76.27 1559
    MISC: precision: 75.51%; recall: 68.87%; FB1: 72.04 841
    ORG: precision: 57.02%; recall: 53.02%; FB1: 54.95 1247
    PER: precision: 39.75%; recall: 44.52%; FB1: 42.00 2063

```

```

[1191]: # 3 more epochs
for epoch in range(3):
    model.train()
    total_loss = 0.0

    # Iterate over your training data in batches
    for batch_idx, (inputs, targets) in enumerate(train_loader):
        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass + src_padding_mask
        src_padding_mask = (inputs == 0).float()
        outputs = model(inputs, src_padding_mask= src_padding_mask)

        # Flatten the outputs and targets for the loss calculation
        outputs = outputs.view(-1, 9)
        targets = targets.view(-1)

        # Calculate the loss
        loss = criterion(outputs, targets)

```

```

    # Backpropagation
    loss.backward()
    optimizer.step()

    total_loss += loss.item()

    # Print the average loss for this epoch
    avg_loss = total_loss / len(train_loader)
    print(f"Epoch [{epoch+22}/{25}] - Loss: {avg_loss:.4f}")

    print('validation error: ')
    precision, recall, f1 = eval(model, val_loader)

```

```

Epoch [22/25] - Loss: 0.1733
validation error:
processed 51362 tokens with 5942 phrases; found: 5506 phrases; correct: 3568.
accuracy: 59.12%; (non-0)
accuracy: 92.24%; precision: 64.80%; recall: 60.05%; FB1: 62.33
      LOC: precision: 83.40%; recall: 69.73%; FB1: 75.96 1536
      MISC: precision: 76.79%; recall: 71.04%; FB1: 73.80 853
      ORG: precision: 57.69%; recall: 53.69%; FB1: 55.62 1248
      PER: precision: 48.80%; recall: 49.51%; FB1: 49.15 1869
Epoch [23/25] - Loss: 0.1700
validation error:
processed 51362 tokens with 5942 phrases; found: 5337 phrases; correct: 3540.
accuracy: 56.79%; (non-0)
accuracy: 92.03%; precision: 66.33%; recall: 59.58%; FB1: 62.77
      LOC: precision: 84.71%; recall: 69.95%; FB1: 76.62 1517
      MISC: precision: 76.77%; recall: 70.61%; FB1: 73.56 848
      ORG: precision: 55.55%; recall: 57.87%; FB1: 56.68 1397
      PER: precision: 52.57%; recall: 44.95%; FB1: 48.46 1575
Epoch [24/25] - Loss: 0.1684
validation error:
processed 51362 tokens with 5942 phrases; found: 5823 phrases; correct: 3586.
accuracy: 58.87%; (non-0)
accuracy: 91.86%; precision: 61.58%; recall: 60.35%; FB1: 60.96
      LOC: precision: 84.38%; recall: 69.41%; FB1: 76.16 1511
      MISC: precision: 76.08%; recall: 68.66%; FB1: 72.18 832
      ORG: precision: 54.89%; recall: 56.08%; FB1: 55.48 1370
      PER: precision: 43.89%; recall: 50.27%; FB1: 46.86 2110

```

```
[1192]: print(f"Validation: precision = {precision}, recall = {recall}, f1 = {f1}")
```

```

Validation: precision = 61.58337626652928, recall = 60.35005048805117, f1 =
60.96047598810029

```

```
[1193]: # SAVE THE MODEL
torch.save(model.state_dict(), 'task3.pth')
```

```
[1186]: #evaluation
def eval(model, loader):
    model.eval()
    all_preds, all_labels = [], []
    with torch.no_grad():
        for batch in loader:
            inputs, targets = batch
            #get rid of paddings on targets
            label_unpad = targets
            mask = label_unpad != 9
            label_unpad = label_unpad[mask]
            src_padding_mask = (inputs == 0).float()
            # print('size match:', src_padding_mask.size() == inputs.size())
            outputs = model(inputs,src_padding_mask=src_padding_mask)
            _, preds = torch.max(outputs, -1)
            #get rid of paddings on pred
            preds = preds[mask]

            preds_converted = [reversed_ner_tags[elem.item()] for elem in preds]
            targets_converted = [reversed_ner_tags[elem.item()] for elem in
↪label_unpad]
            all_preds.extend(preds_converted)
            all_labels.extend(targets_converted)
            # all_preds = list(chain.from_iterable(all_preds))
            # all_labels = list(chain.from_iterable(all_labels))
            # all_labels = torch.cat(all_labels)
            # all_preds = itertools.chain(*all_preds)
            # all_labels =itertools.chain(*all_labels)
            result = evaluate(all_labels, all_preds,verbose=True)
            precision, recall, f1 = result[0], result[1],result[2]
        return precision, recall, f1
```

```
[1194]: print('Test: ')
precision, recall, f1 = eval(model, test_loader)
```

Test:

processed 46435 tokens with 5648 phrases; found: 5360 phrases; correct: 2830.

accuracy: 50.48%; (non-0)

accuracy: 89.56%; precision: 52.80%; recall: 50.11%; FB1: 51.42

LOC: precision: 79.85%; recall: 64.15%; FB1: 71.14 1340

MISC: precision: 67.76%; recall: 61.97%; FB1: 64.73 642

ORG: precision: 49.84%; recall: 45.88%; FB1: 47.77 1529

PER: precision: 30.45%; recall: 34.82%; FB1: 32.49 1849

```
[1195]: print(f"Test: precision = {precision}, recall = {recall}, f1 = {f1}")
```

Test: precision = 52.79850746268657, recall = 50.106232294617556, f1 = 51.417151162790695

0.2.1 Solution to task 3

1. Hyperparameters:
 - `embedding_dim = 100`
 - `hidden_dim = 256`
 - `output_dim = 128`
 - `num_layers = 1`
 - `dropout = 0.33`
 - `optimizer learning rate = 0.001`
 - `ignore_index = 9`
 - `batch_size = 32`
2. Solution: Same as task, we still use `input_ids` as the input. The dataloader will pad 0 to input and 9 to `ner_tags`. Next, the first layer of the model is an embedding layer, which convert each token into 128-d vector. The positional encoder is a self-attention layer which will generate a sequence as output. For `src_padding_mask`, it will identify all the padded values and get rid of their impact. Next, the batches will be thrown to the transformer encoder and a FFN to predict the results.
3. Questions and answers:
 - What is the precision, recall, and F1 score on the validation data?
 - `precision = 61.58337626652928`, `recall = 60.35005048805117`, `f1 = 60.96047598810029`
 - What are the precision, recall, and F1 score on the test data?
 - `precision = 52.79850746268657`, `recall = 50.106232294617556`, `f1 = 51.417151162790695`-
What is the reason behind the poor performance of the transformer?
 - At first, the transformer typically require big amount of data. Since the `word2idx` is too small, it cannot generalize well onto the new data. Secondly, the other problem of the small dataset is that the model will probably overfit.